

Assignment:

Hdfs task-1

1. What is the Namenode's URI and which file is it configured in?

Namenode's URI is 9000, it is configured in `core-site.xml`

2. Where on a local file system will Namenode store its image and which file is it configured in?

`Hadoop-env.sh`

3. Where on a local file system will Datanode store its blocks and which file is it configured in?

`Hadoop-env.sh`

4. What is the block replication and which file is it configured in?

Block replication stores large amount of data in sequence of blocks around cluster.

Perform:

Start HDFS and verify that it's running

```
Select hduser@LAPTOP-CBCQ82TI: ~/hadoop/hadoop-3.3.0
hduser@LAPTOP-CBCQ82TI:~$ cd ~/hadoop/hadoop-3.3.0/
hduser@LAPTOP-CBCQ82TI:~/hadoop/hadoop-3.3.0$ ls
LICENSE-binary LICENSE.txt NOTICE-binary NOTICE.txt README.txt bin etc include lib libexec licenses-binary logs sbin share
hduser@LAPTOP-CBCQ82TI:~/hadoop/hadoop-3.3.0$ sudo service ssh status
[sudo] password for hduser:
sshd is not running
hduser@LAPTOP-CBCQ82TI:~/hadoop/hadoop-3.3.0$ sudo service ssh restart
* Restarting OpenBSD Secure Shell server sshd
hduser@LAPTOP-CBCQ82TI:~/hadoop/hadoop-3.3.0$ sudo service ssh status
* sshd is running
hduser@LAPTOP-CBCQ82TI:~/hadoop/hadoop-3.3.0$ ls
LICENSE-binary LICENSE.txt NOTICE-binary NOTICE.txt README.txt bin etc include lib libexec licenses-binary logs sbin share
hduser@LAPTOP-CBCQ82TI:~/hadoop/hadoop-3.3.0$ ls bin/
container-executor hadoop hadoop.cmd hdfs hdfs.cmd mapred mapred.cmd oom-listener test-container-executor yarn yarn.cmd
hduser@LAPTOP-CBCQ82TI:~/hadoop/hadoop-3.3.0$ ls sbin/
 federationStateStore https.sh start-all.cmd start-dfs.sh stop-all.cmd stop-dfs.sh workers.sh
 distribute-exclude.sh kms.sh start-all.sh start-secure-dns.sh stop-all.sh stop-secure-dns.sh yarn-daemon.sh
 hadoop-daemon.sh mr-jobhistory-daemon.sh start-balancer.sh start-yarn.cmd stop-balancer.sh stop-yarn.cmd yarn-daemons.sh
 hadoop-daemons.sh refresh-namenodes.sh start-dfs.cmd start-yarn.sh stop-dfs.cmd stop-yarn.sh
hduser@LAPTOP-CBCQ82TI:~/hadoop/hadoop-3.3.0$ sbin/start-dfs.sh
Starting namenodes on [localhost]
Starting datanodes
Starting secondary namenodes [LAPTOP-CBCQ82TI]
hduser@LAPTOP-CBCQ82TI:~/hadoop/hadoop-3.3.0$ jps
952 Jps
489 DataNode
316 NameNode
734 SecondaryNameNode
hduser@LAPTOP-CBCQ82TI:~/hadoop/hadoop-3.3.0$ sbin/start-yarn.sh
Starting resourcemanager
Starting nodemanagers
hduser@LAPTOP-CBCQ82TI:~/hadoop/hadoop-3.3.0$ jps
1283 NodeManager
1398 Jps
1095 ResourceManager
489 DataNode
316 NameNode
734 SecondaryNameNode
hduser@LAPTOP-CBCQ82TI:~/hadoop/hadoop-3.3.0$
```

Check the entire filesystem for inconsistencies/problems

```

hduser@LAPTOP-C8CQB2TI:~/hadoop/hadoop-3.3.0$ bin/hdfs fsck /exercise1
Connecting to namenode via http://localhost:9870/fsck?ugi=hduser&path=%2Fexercise1
FSCK started by hduser (auth:SIMPLE) from /127.0.0.1 for path /exercise1 at Wed Jun 02 17:26:01 IST 2021

Status: HEALTHY
  Number of data-nodes: 1
  Number of racks:      1
  Total dirs:           1
  Total symlinks:       0

Replicated Blocks:
  Total size: 743 B
  Total files: 1
  Total blocks (validated): 1 (avg. block size 743 B)
  Minimally replicated blocks: 1 (100.0 %)
  Over-replicated blocks: 0 (0.0 %)
  Under-replicated blocks: 0 (0.0 %)
  Mis-replicated blocks: 0 (0.0 %)
  Default replication factor: 1
  Average block replication: 1.0
  Missing blocks: 0
  Corrupt blocks: 0
  Missing replicas: 0 (0.0 %)
  Blocks queued for replication: 0

Erasure Coded Block Groups:
  Total size: 0 B
  Total files: 0
  Total block groups (validated): 0
  Minimally erasure-coded block groups: 0
  Over-erasure-coded block groups: 0
  Under-erasure-coded block groups: 0
  Unsatisfactory placement block groups: 0
  Average block group size: 0.0
  Missing block groups: 0
  Corrupt block groups: 0
  Missing internal blocks: 0
  Blocks queued for replication: 0
FSCK ended at Wed Jun 02 17:26:01 IST 2021 in 27 milliseconds

The filesystem under path '/exercise1' is HEALTHY
hduser@LAPTOP-C8CQB2TI:~/hadoop/hadoop-3.3.0$

```

Copy deckofcards.txt back to local file system and name it deckofcards.copy.txt

```

hduser@LAPTOP-C8CQB2TI:~/hadoop/hadoop-3.3.0$ bin/hadoop fs -get /exercise1/deckofcards.txt /home/hduser/sunil
hduser@LAPTOP-C8CQB2TI:~/hadoop/hadoop-3.3.0$

```

Copy deckofcards.txt to deckofcardsCopy.txt

```

hduser@LAPTOP-C8CQ82T1: ~/hadoop/hadoop-1.3.0$ bin/hadoop distcp /exercise1/deckofcards.txt /new_fol/copydeckofcards.txt
2021-06-02 17:02:34,431 INFO tools.DistCp: Input Options: DistCpOptions(atomicCommit=false, deleteMissing=false, ignoreFailures=false, overwrite=false,
, append=false, useDiff=false, useDiff=false, fromSnapshot=null, toSnapshot=null, skipCRC=false, blocking=true, numListStatusThreads=0, maxMaps=20, mapBandwidth=0.0, c
pyStrategy="uniformsize", preserveStatus=[BLOCKSIZE], atomicWorkPath=null, logPath=null, sourceFileListing=null, sourcePaths=[/exercise1/deckofcards.txt], targetPath=/
new_fol/copydeckofcards.txt, filtersFile="null", blocksPerChunk=0, copyBufferSize=8192, verboseLog=false, directWrite=false}, sourcePaths=[/exercise1/deckofcards.txt],
targetPathExists=false, preserveRawXattrs=false
2021-06-02 17:02:34,958 INFO client.DefaultNoHARMFailoverProxyProvider: Connecting to ResourceManager at /0.0.0.0:8032
2021-06-02 17:02:40,236 INFO tools.SimpleCopyListing: Paths (files+dirs) cnt = 1; dirCnt = 0
2021-06-02 17:02:40,237 INFO tools.SimpleCopyListing: Build file listing completed.
2021-06-02 17:02:40,240 INFO Configuration.deprecation: io.sort.mb is deprecated. Instead, use mapreduce.task.io.sort.mb
2021-06-02 17:02:40,241 INFO Configuration.deprecation: io.sort.factor is deprecated. Instead, use mapreduce.task.io.sort.factor
2021-06-02 17:02:40,981 INFO tools.DistCp: Number of paths in the copy list: 1
2021-06-02 17:02:41,151 INFO tools.DistCp: Number of paths in the copy list: 1
2021-06-02 17:02:41,183 INFO client.DefaultNoHARMFailoverProxyProvider: Connecting to ResourceManager at /0.0.0.0:8032
2021-06-02 17:02:43,397 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /tmp/hadoop-yarn/staging/hduser/.staging/job_1622630291232_0001
2021-06-02 17:02:44,122 INFO mapreduce.JobSubmitter: number of splits:1
2021-06-02 17:02:46,471 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1622630291232_0001
2021-06-02 17:02:46,474 INFO mapreduce.JobSubmitter: Executing with tokens: []
2021-06-02 17:02:47,657 INFO conf.Configuration: resource-types.xml not found
2021-06-02 17:02:47,658 INFO resource.ResourceUtils: Unable to find 'resource-types.xml'.
2021-06-02 17:02:51,937 INFO Impl.YarnClientImpl: Submitted application application_1622630291232_0001
2021-06-02 17:02:51,483 INFO mapreduce.Job: The url to track the job: http://LAPTOP-C8CQ82T1.localdomain:8088/proxy/application_1622630291232_0001/
2021-06-02 17:02:51,484 INFO tools.DistCp: DistCp job-id: job_1622630291232_0001
2021-06-02 17:02:51,584 INFO mapreduce.Job: Running job: job_1622630291232_0001
2021-06-02 17:04:32,235 INFO mapreduce.Job: Job job_1622630291232_0001 running in uber mode : false
2021-06-02 17:04:32,283 INFO mapreduce.Job: map 0% reduce 0%
2021-06-02 17:05:28,217 INFO mapreduce.Job: map 100% reduce 0%
2021-06-02 17:05:36,410 INFO mapreduce.Job: Job job_1622630291232_0001 completed successfully
2021-06-02 17:05:36,769 INFO mapreduce.Job: Counters: 37
  File System Counters
    FILE: Number of bytes read=0
    FILE: Number of bytes written=267605
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=1130
    HDFS: Number of bytes written=743
    HDFS: Number of read operations=16
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=5
    HDFS: Number of bytes read erasure-coded=0
  Job Counters
    Launched map tasks=1
    Other local map tasks=1
    Total time spent by all maps in occupied slots (ms)=49089
    Total time spent by all reduces in occupied slots (ms)=0
    Total time spent by all map tasks (ms)=49089
    Total vcore-millisecons taken by all map tasks=49089
    Total megabyte-millisecons taken by all map tasks=58267136
  Map-Reduce Framework
    Map input records=1
    Map output records=0
    Input split bytes=136
    Spilled Records=0
    Failed Shuffles=0
    Merged Map outputs=0
    GC time elapsed (ms)=94
    CPU time spent (ms)=2808
    Physical memory (bytes) snapshot=165036032
    Virtual memory (bytes) snapshot=1674669531136
    Total committed heap usage (bytes)=69206016
    Peak Map Physical memory (bytes)=165036032
    Peak Map Virtual memory (bytes)=1674669531136
  File Input Format Counters
    Bytes Read=251
  File Output Format Counters
    Bytes Written=0
  DistCp Counters
    Bandwidth in Bytes=371
    Bytes Copied=743
    Bytes Expected=743
    Files Copied=1
hduser@LAPTOP-C8CQ82T1: ~/hadoop/hadoop-1.3.0$ bin/hadoop fs -ls /exercise1
Found 1 items
-rw-r--r-- 1 hduser supergroup 743 2021-06-02 16:38 /exercise1/deckofcards.txt
hduser@LAPTOP-C8CQ82T1: ~/hadoop/hadoop-1.3.0$ bin/hadoop fs -ls /new_fol
Found 2 items
-rw-r--r-- 1 hduser supergroup 28620 2021-05-25 19:13 /new_fol/copy_titanic.csv
-rw-r--r-- 1 hduser supergroup 743 2021-06-02 17:05 /new_fol/copydeckofcards.txt
hduser@LAPTOP-C8CQ82T1: ~/hadoop/hadoop-1.3.0$

```

Upload GitHub repo sample_data/deckofcards.txt to HDFS under /exercise1 directory

```

hduser@LAPTOP-C8CQB2TI:~/hadoop/hadoop-3.3.0$ bin/hdfs dfs -copyFromLocal /home/hduser/deckofcards.txt /exercise1
hduser@LAPTOP-C8CQB2TI:~/hadoop/hadoop-3.3.0$ bin/hdfs dfs -ls /exercise1
Found 1 items
-rw-r--r-- 1 hduser supergroup 743 2021-06-02 16:38 /exercise1/deckofcards.txt
hduser@LAPTOP-C8CQB2TI:~/hadoop/hadoop-3.3.0$

```

Create a new directory /exercise1 on HDFS

```

hduser@LAPTOP-C8CQB2TI:~/hadoop/hadoop-3.3.0$ bin/hdfs dfs -ls /sunil
Found 1 items
-rw-r--r-- 1 hduser supergroup 28629 2021-05-25 19:13 /sunil/titanic_test.csv
hduser@LAPTOP-C8CQB2TI:~/hadoop/hadoop-3.3.0$ bin/hdfs dfs -mkdir /exercise1
hduser@LAPTOP-C8CQB2TI:~/hadoop/hadoop-3.3.0$ bin/hdfs dfs -ls /
Found 3 items
drwxr-xr-x - hduser supergroup 0 2021-06-02 16:18 /exercise1
drwxr-xr-x - hduser supergroup 0 2021-05-27 15:57 /new_fol
drwxr-xr-x - hduser supergroup 0 2021-05-25 19:13 /sunil
hduser@LAPTOP-C8CQB2TI:~/hadoop/hadoop-3.3.0$

```

Delete the /exercise1 directory from HDFS

```

hduser@LAPTOP-C8CQB2TI:~/hadoop/hadoop-3.3.0$ bin/hdfs dfs -ls /
Found 4 items
drwxr-xr-x - hduser supergroup 0 2021-06-02 17:27 /exercise1
drwxr-xr-x - hduser supergroup 0 2021-06-02 17:05 /new_fol
drwxr-xr-x - hduser supergroup 0 2021-06-02 16:53 /sunil
drwx----- - hduser supergroup 0 2021-06-02 17:02 /tmp
hduser@LAPTOP-C8CQB2TI:~/hadoop/hadoop-3.3.0$ bin/hdfs dfs -rm -r /exercise1
Deleted /exercise1
hduser@LAPTOP-C8CQB2TI:~/hadoop/hadoop-3.3.0$ bin/hdfs dfs -ls /
Found 3 items
drwxr-xr-x - hduser supergroup 0 2021-06-02 17:05 /new_fol
drwxr-xr-x - hduser supergroup 0 2021-06-02 16:53 /sunil
drwx----- - hduser supergroup 0 2021-06-02 17:02 /tmp
hduser@LAPTOP-C8CQB2TI:~/hadoop/hadoop-3.3.0$

```

Delete deckofcards.txt from HDFS


```
hduser@LAPTOP-C8CQB2TI:~/hadoop/hadoop-3.3.4$ bin/hdfs dfs -ls /exercise1
Found 1 items
-rw-r--r-- 1 hduser supergroup 743 2021-06-02 16:38 /exercise1/deckofcards.txt
hduser@LAPTOP-C8CQB2TI:~/hadoop/hadoop-3.3.4$ bin/hdfs dfs -rm /exercise1/deckofcards.txt
Deleted /exercise1/deckofcards.txt
hduser@LAPTOP-C8CQB2TI:~/hadoop/hadoop-3.3.4$ bin/hdfs dfs -ls /exercise1
hduser@LAPTOP-C8CQB2TI:~/hadoop/hadoop-3.3.4$
```

Print the first 25 lines to the screen from deckofcards.txt on HDFS

```
hduser@LAPTOP-C8CQB2TI:~/hadoop/hadoop-3.3.4$ bin/hdfs dfs -cat /exercise1/deckofcards.txt | head -n 25
BLACK|SPADE|2
BLACK|SPADE|3
BLACK|SPADE|4
BLACK|SPADE|5
BLACK|SPADE|6
BLACK|SPADE|7
BLACK|SPADE|8
BLACK|SPADE|9
BLACK|SPADE|10
BLACK|SPADE|J
BLACK|SPADE|Q
BLACK|SPADE|K
BLACK|SPADE|A
BLACK|CLUB|2
BLACK|CLUB|3
BLACK|CLUB|4
BLACK|CLUB|5
BLACK|CLUB|6
BLACK|CLUB|7
BLACK|CLUB|8
BLACK|CLUB|9
BLACK|CLUB|10
BLACK|CLUB|J
BLACK|CLUB|Q
BLACK|CLUB|K
hduser@LAPTOP-C8CQB2TI:~/hadoop/hadoop-3.3.4$
```

View the content of the /exercise1 directory

```
hduser@LAPTOP-C8CQB2TI:~/hadoop/hadoop-3.3.4$ bin/hdfs dfs -cat /exercise1/deckofcards.txt
BLACK|SPADE|2
BLACK|SPADE|3
BLACK|SPADE|4
BLACK|SPADE|5
BLACK|SPADE|6
BLACK|SPADE|7
BLACK|SPADE|8
BLACK|SPADE|9
BLACK|SPADE|10
BLACK|SPADE|J
BLACK|SPADE|Q
BLACK|SPADE|K
BLACK|SPADE|A
BLACK|CLUB|2
BLACK|CLUB|3
BLACK|CLUB|4
BLACK|CLUB|5
BLACK|CLUB|6
BLACK|CLUB|7
BLACK|CLUB|8
BLACK|CLUB|9
BLACK|CLUB|10
BLACK|CLUB|J
BLACK|CLUB|Q
BLACK|CLUB|K
BLACK|CLUB|A
RED|DIAMOND|2
RED|DIAMOND|3
RED|DIAMOND|4
```